

Introduction to Nonlinear Solvers Using PETSc/TAO

Presented to
ATPESC 2024 Participants

Toby Isaac (from slides by Richard Tran Mills)
Argonne National Laboratory
August 6, 2024



ATPESC Numerical Software Track



Iterative Solvers for Nonlinear Systems

Systems of nonlinear equations

$$F(x) = b \quad \text{where} \quad F : \mathbb{R}^N \rightarrow \mathbb{R}^N \quad (1)$$

arise in countless settings in computational science.

Direct methods for general nonlinear systems do not exist. Iterative methods are required!

Nonlinear Richardson (simple) iteration:

$$x_{k+1} = x_k + \lambda(b - F(x_k)) \quad (2)$$

This has linear convergence at best: $\|e_{k+1}\| \leq C\|e_k\|$

Nonlinear Krylov methods

Nonlinear CG - Mimic CG to force each new search direction to be orthogonal to previous directions.

Nonlinear GMRES (Anderson mixing) - minimize $\|F(x_{k+1}) - b\|$ by using x_{k+1} as a linear combination of previous solutions and solving a linear least squares problem.

These have superlinear convergence at best: $\|e_{k+1}\| \leq C\|e_k\|^{\alpha \geq 1}$

Newton's Method: The workhorse of nonlinear solvers

- Standard form of a nonlinear system

$$F(u) = 0$$

- Iteration

$$\text{Solve: } J(u)w = -F(u)$$

$$\text{Update: } u^+ \leftarrow u + w$$


Where the Jacobian $J(u) = F'(u) = \frac{\partial F(u)}{\partial u}$.

- Quadratically convergent near a root: $|u^{n+1} - u^*| \in \mathcal{O}(|u^n - u^*|^2)$

Example (Nonlinear Poisson)

$$F(u) = 0 \quad \sim \quad -\nabla \cdot [(1 + u^2)\nabla u] - f = 0$$

$$J(u)w \quad \sim \quad -\nabla \cdot [(1 + u^2)\nabla w + 2uw\nabla u]$$



High-Performance Nonlinear Solvers Are Available Via



PETSc: Portable, Extensible Toolkit for Scientific Computation



SUNDIALS: SUite of Nonlinear and Differential/ALgebraic Equation Solvers



Trilinos: NOX (Nonlinear Object-Oriented Solutions) and LOCA (Library of Continuation Algorithms) packages.

This presentation focuses on PETSc.

Optimization

Time Integrators

Nonlinear Algebraic Solvers

Krylov Subspace Solvers

Preconditioners

**Domain-
Specific
Interfaces**

Networks

Quadtree / Octree

Unstructured Mesh

Structured Mesh

Vectors

Index Sets

Matrices

**Computation & Communication
Kernels**


■ Easy customization and composability of solvers at runtime

- Enables optimality via flexible combinations of physics, algorithmics, architectures
- Try new algorithms by composing new/existing algorithms (multilevel, domain decomposition, splitting, etc.)

■ Portability & performance

- Largest DOE machines, also clusters, laptops
- Thousands of users worldwide

Argonne
NATIONAL LABORATORY



PETSc provides the backbone of diverse scientific applications.

clockwise from upper left: hydrology, cardiology, fusion, multiphase steel, relativistic matter, ice sheet modeling



<https://www.mcs.anl.gov/petsc>

PETSc - the Portable, Extensible Toolkit for Scientific computation

Application Codes

Higher-Level Libraries and Frameworks

**TS**

Time Steppers

Pseudo-Transient, Runge-Kutta, IMEX, SSP, ...
Local and Global Error Estimators
Adaptive Timestepping
Event Handling
Sensitivity via Adjoints

SNES

Nonlinear Solvers

Newton Linesearch Successive Substitution
Newton Trust Region Nonlinear CG
BFGS (Quasi-Newton) Active Set VI
Nonlinear Gauss-Seidel

KSP

Linear Solvers

CG, GMRES, BiCGStab, FGMRES, ...
Pipelined Krylov Methods
Hierarchical Krylov Methods

PC

Preconditioners

ILU/ICG
Additive Schwarz
Fieldsplit (Block Preconditioners)
PCMG (Geometric Multigrid)
GAMG (Algebraic Multigrid)

Vec

Vectors

IS

Index Sets

Mat

Linear Operators

AIJ (Compressed Sparse Row)
SAIJ (Symmetric)
BAIJ (Blocked)
Dense
GPU Matrices



Optimization Solvers

PDE-Constrained
Adjoint-Based
Derivative-Free

Levenberg-Marquardt
Newton's Method
Interior Point Methods



Domain Management



DMDA
Regular Grids



DMplex
Unstructured Meshes



DMForest
Forest-of-trees AMR



DMStag
Staggered Grids



DMNetwork
Networks



DMSwarm
Particles

PetscSF

Parallel Communication

SLEPc
Eigensolvers

Communication and Computational Kernels

MPI

BLAS/LAPACK

Kokkos

CUDA

...



PETSc is a platform for experimentation!

Optimal solvers must consider the interplay among physics, algorithms, and architectures. This need motivates several elements of the PETSc design philosophy:

Algebraic solvers must be

- Composable: Separately developed solvers should be easy to combine, by non-experts, to form a more powerful solver.
- Hierarchical: Outer solvers may iterate over all variables for a global problem, while inner solvers handle smaller subsets of physics, smaller physical subdomains, or coarser meshes.
- Nested: Outer solvers call nested inner solvers
- Extensible: Easily customized or extended by users

To facilitate experimentation, many solver configurations can be set at runtime; no need to recompile. It is also possible to control aspects of data layout, choice of computational kernels, architecture to execute on, and more via the runtime options database.

The philosophy of supporting extensive runtime experimentation guides many of the design choices made in PETSc (including the GPU usage model, which we'll briefly look at if time permits).

In practice, incurring the expense of an exact solve for the Newton step is often not desirable:

Inexact Newton methods find an approximate Newton direction Δx_k that satisfies

$$\|F'(x_k)\Delta x_k + F(x_k)\| \leq \eta \|F(x_k)\| \quad (3)$$

for a forcing term $\eta \in [0, 1)$ (static or chosen adaptively via Eisenstat-Walker method, `-snes_ksp_ew`).

Newton-Krylov methods, which use Krylov subspace projection methods as the “inner”, linear iterative solver, are a robust and widely-used variant.

PETSc provides a wide range of Krylov methods and linear preconditioners that can be accessed via runtime options (`-ksp_type <ksp_method> -pc_type <pc_method>`).

Newton has quadratic convergence *only when the iterate is sufficiently close to the solution*. Far from the solution, the computed Newton step is often too large in magnitude.

In practice, some globalization strategy is often needed to expand the domain of convergence. PETSc offers several options; most common (and default) is backtracking line search.

Backtracking line search

- Replaces the full Newton step s with some scalar multiple: $x_{k+1} = x_k + \lambda_k s$, $\lambda > 0$
- Introduce merit function $\phi(x) = \frac{1}{2} \|F(x)\|_2^2$, (approximately) find $\min_{\lambda > 0} \phi(x_k + \lambda s)$
- Accurate minimization not worth the expense; simply ensure sufficient decrease:

$$\phi(x_k + \lambda s) \leq \phi(x_k) + \alpha \lambda s^T \nabla \phi(x_k) \quad (4)$$

(α is a user-tunable parameter; defaults to $1e-4$)

- Builds polynomial model for $\phi(x_k + \lambda s)$ (default is cubic; change via `-snes_linesearch_order <n>`).

The PETSc SNES (Scalable Nonlinear Equation Solvers) interface is based upon callback functions

- `FormFunction()`, set by `SNESSetFunction()`
- `FormJacobian()`, set by `SNESSetJacobian()`

When PETSc needs to evaluate the nonlinear residual $F(x)$,

- Solver calls the **user's** function
- User function gets application state through the `ctx` variable
 - PETSc *never* sees application data

SNES Function

The user provided function which calculates the nonlinear residual has signature

```
PetscErrorCode (*func)(SNES snes,Vec x,Vec r,void *ctx)
```

x: The current solution

r: The residual

ctx: The user context passed to SNESSetFunction()

- Use this to pass application information, e.g. physical constants

The user provided function that calculates the Jacobian has signature

```
PetscErrorCode (*func)(SNES snes, Vec x, Mat J,  
                        Mat Jpre, void *ctx)
```

x: The current solution

J: The Jacobian

Jpre: The Jacobian preconditioning matrix (possibly J itself)

ctx: The user context passed to SNESSetFunction()


- Use this to pass application information, e.g. physical constants

Alternatively, you can use


- a builtin sparse finite difference approximation (“coloring”)

SNES Example ex19, Steady-State Nonisothermal Lid-Driven Cavity


Solution Components



velocity: u



velocity: v



vorticity:



temperature: T

$$-\Delta U - \partial_y \Omega = 0$$

$$-\Delta V + \partial_x \Omega = 0$$

$$-\Delta \Omega + \nabla \cdot ([U\Omega, V\Omega]) - \text{Gr} \partial_x T = 0$$

$$-\Delta T + \text{Pr} \nabla \cdot ([UT, VT]) = 0$$

- 2D square domain with a moving lid; nonisothermal (temperature T)
- Flow driven by lid motion and buoyancy effects
- Velocity(U, V)-vorticity(Ω) formulation
- Finite difference discretization
- Logically regular grid
 - Parallelized with DMDA
- Analytical Jacobian not provided; Calculated by finite-differences (using coloring)
- Contributed by David Keyes
- Run as part of `make check`

SNES Example

Driven Cavity Application Context

```
/*  
    User-defined data structures  
*/  
  
/* Collocated at each node */  
typedef struct {  
    PetscScalar u,v,omega,temp;  
} Field;  
  
typedef struct {  
    PetscReal lidvelocity,prandtl,grashof; /* physical parameters */  
    PetscBool draw_contours; /* flag - 1 indicates drawing countours */  
} AppCtx;
```

SNES Function with local evaluation

```
PetscErrorCode FormFunctionLocal(DMDALocalInfo *info,Field **x,Field **f,void *ptr)
{
    ...
    xints = info->xs; xinte = info->xs+info->xm; yints = info->ys; yinte = info->ys+info->ym;
    /* Handle boundaries ... */
    /* Compute over the interior points */

    for (j=yints; j<yinte; j++) {
        for (i=xints; i<xinte; i++) {
            /* convective coefficients for upwinding ... */
            /* U velocity */
            u          = x[j][i].u;
            uxx         = (2.0*u - x[j][i-1].u - x[j][i+1].u)*hydhx;
            uyy         = (2.0*u - x[j-1][i].u - x[j+1][i].u)*hxdhy;
            f[j][i].u    = uxx + uyy - .5*(x[j+1][i].omega-x[j-1][i].omega)*hx;
            /* V velocity, Omega ... */
            /* Temperature */
            u          = x[j][i].temp;
            uxx         = (2.0*u - x[j][i-1].temp - x[j][i+1].temp)*hydhx;
            uyy         = (2.0*u - x[j-1][i].temp - x[j+1][i].temp)*hxdhy;
            f[j][i].temp = uxx + uyy + prandtl
                * ( (vxp*(u - x[j][i-1].temp) + vxm*(x[j][i+1].temp - u)) * hy
                    + (vyp*(u - x[j-1][i].temp) + vym*(x[j+1][i].temp - u)) * hx);
        }
    }
}
```

\$PETSC_DIR/src/snes/tutorials/ex19.c

Hands-on: Running the driven cavity

Run SNES ex19 with a single MPI rank (see full instructions for all hands-on exercises [here](#)):

```
./ex19 -snes_monitor -snes_converged_reason -da_grid_x 16 -da_grid_y 16 -da_refine  
2 -lidvelocity 100 -grashof 1e2
```


Hands-on: Running the driven cavity

Run SNES ex19 with a single MPI rank (see full instructions for all hands-on exercises [here](#)):

```
./ex19 -snes_monitor -snes_converged_reason -da_grid_x 16 -da_grid_y 16 -da_refine  
2 -lidvelocity 100 -grashof 1e2
```

Explanation of the above command-line options:

- `-snes_monitor`: Show progress of the SNES solver
- `-snes_converged_reason`: Print reason for SNES convergence or divergence
- `-da_grid_x 16`: Set initial grid points in x direction to 16
- `-da_grid_y 16`: Set initial grid points in y direction to 16
- `-da_refine 2`: Refine the initial grid 2 times before creation
- `-lidvelocity 100`: Set dimensionless velocity of lid to 100
- `-grashof 1e2`: Set Grashof number to 1e2

An element of the PETSc design philosophy is extensive runtime customizability; Use `-help` to enumerate and explain the various command-line options available.

Hands-on: Running the driven cavity

Run SNES ex19 with a single MPI rank (see full instructions for all hands-on exercises [here](#)):

```
./ex19 -snes_monitor -snes_converged_reason -da_grid_x 16 -da_grid_y 16 -da_refine  
2 -lidvelocity 100 -grashof 1e2
```

```
lid velocity = 100., prandtl # = 1., grashof # = 100.  
0 SNES Function norm 7.681163231938e+02  
1 SNES Function norm 6.582880149343e+02  
2 SNES Function norm 5.294044874550e+02  
3 SNES Function norm 3.775102116141e+02  
4 SNES Function norm 3.047226778615e+02  
5 SNES Function norm 2.599983722908e+00  
6 SNES Function norm 9.427314747057e-03  
7 SNES Function norm 5.212213461756e-08  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 7  
Number of SNES iterations = 7
```

What is the SNES solver actually doing? Add `-snes_view` to see

```
SNES Object: 1 MPI processes
  type: newtonls
  maximum iterations=50, maximum function evaluations=10000
  tolerances: relative=1e-08, absolute=1e-50, solution=1e-08
  total number of linear solver iterations=835
  total number of function evaluations=11
  norm schedule ALWAYS
  Jacobian is built using colored finite differences on a DMDA
  SNESLineSearch Object: 1 MPI processes
    type: bt
      interpolation: cubic
      alpha=1.000000e-04
      maxstep=1.000000e+08, minlambda=1.000000e-12
      tolerances: relative=1.000000e-08, absolute=1.000000e-15, lambda=1.000000e-08
      maximum iterations=40
  KSP Object: 1 MPI processes
    type: gmres
      restart=30, using Classical (unmodified) Gram-Schmidt Orthogonalization with no iterative refinement
      happy breakdown tolerance 1e-30
      maximum iterations=10000, initial guess is zero
      tolerances: relative=1e-05, absolute=1e-50, divergence=10000.
      left preconditioning
      using PRECONDITIONED norm type for convergence test
  PC Object: 1 MPI processes
    type: ilu
      out-of-place factorization
      0 levels of fill
      tolerance for zero pivot 2.22045e-14
      matrix ordering: natural
      factor fill ratio given 1., needed 1.
      Factored matrix follows:
        Mat Object: 1 MPI processes
          type: seqaij
          rows=14884, cols=14884, bs=4
          package used to perform factorization: petsc
          total: nonzeros=293776, allocated nonzeros=293776
```

Managing PETSc options

PETSc offers a very large number of runtime options.

All can be set via command line, but can also be set from input files and shell environment variables.

To facilitate readability, we'll put the command-line arguments common to the remaining hands-on exercises in PETSC_OPTIONS.

```
export PETSC_OPTIONS="-snes_monitor -snes_converged_reason -lidvelocity 100  
-da_grid_x 16 -da_grid_y 16 -ksp_converged_reason -log_view :log.txt"
```

We've added `-ksp_converged_reason` to see how and when linear solver halts.

We've also added `-log_view` to write the PETSc performance logging info to a file.

We don't have time to explain the performance logs; find the overall wall-clock time via

```
grep Time\ \(sec\): log.txt
```

Hands-on: Exact vs. Inexact Newton

PETSc defaults to inexact Newton. To run exact (and check the execution time), do

```
./ex19 -da_refine 2 -grashof 1e2 -pc_type lu  
grep Time\ \ (sec\): log.txt
```

Now run inexact Newton and vary the liner solve tolerance (`-ksp_rtol`).

```
./ex19 -da_refine 2 -grashof 1e2 -ksp_rtol 1e-8  
./ex19 -da_refine 2 -grashof 1e2 -ksp_rtol 1e-5  
./ex19 -da_refine 2 -grashof 1e2 -ksp_rtol 1e-3  
./ex19 -da_refine 2 -grashof 1e2 -ksp_rtol 1e-2  
./ex19 -da_refine 2 -grashof 1e2 -ksp_rtol 1e-1
```

What happens to the SNES iteration count? When does it diverge?

What yields the shortest execution time?

Hands-on: Scaling up grid size and running in parallel

What happens to iteration counts (and execution time) as we scale up the grid size?

For this exercise, run in parallel because experiments may take too long otherwise.

We also use BiCGStab (`-ksp_type bcgs`) because the default GMRES(30) fails for some cases.

Hands-on: Scaling up grid size and running in parallel

What happens to iteration counts (and execution time) as we scale up the grid size?

For this exercise, run in parallel because experiments may take too long otherwise.

We also use BiCGStab (`-ksp_type bcgs`) because the default GMRES(30) fails for some cases.

Run with default preconditioner. What happens to iteration counts and execution time?

```
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -da_refine 2  
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -da_refine 3  
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -da_refine 4
```

Hands-on: Scaling up grid size and running in parallel

What happens to iteration counts (and execution time) as we scale up the grid size?

For this exercise, run in parallel because experiments may take too long otherwise.

We also use BiCGStab (`-ksp_type bcgs`) because the default GMRES(30) fails for some cases.

Run with default preconditioner. What happens to iteration counts and execution time?

```
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -da_refine 2
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -da_refine 3
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -da_refine 4
```

```
$ mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -da_refine 4
lid velocity = 100., prandtl # = 1., grashof # = 100.
0 SNES Function norm 1.545962539057e+03
Linear solve converged due to CONVERGED_RTOL iterations 172
1 SNES Function norm 9.780980584978e+02
Linear solve converged due to CONVERGED_RTOL iterations 128
2 SNES Function norm 6.620854219003e+02
Linear solve converged due to CONVERGED_RTOL iterations 600
3 SNES Function norm 3.219025282761e+00
Linear solve converged due to CONVERGED_RTOL iterations 470
4 SNES Function norm 9.280944447516e-03
Linear solve converged due to CONVERGED_RTOL iterations 467
5 SNES Function norm 1.354460792476e-07
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 5
```


Hands-on: Scaling up grid size and running in parallel

What happens to iteration counts (and execution time) as we scale up the grid size?

For this exercise, run in parallel because experiments may take too long otherwise.

We also use BiCGStab (`-ksp_type bcgs`) because the default GMRES(30) fails for some cases.

Let's try geometric multigrid (defaults to V-cycle) by adding `-pc_type mg`

```
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -pc_type mg -da_refine 2
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -pc_type mg -da_refine 3
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -pc_type mg -da_refine 4
```

Hands-on: Scaling up grid size and running in parallel

What happens to iteration counts (and execution time) as we scale up the grid size?

For this exercise, run in parallel because experiments may take too long otherwise.

We also use BiCGStab (`-ksp_type bcgs`) because the default GMRES(30) fails for some cases.

Let's try geometric multigrid (defaults to V-cycle) by adding `-pc_type mg`

```
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -pc_type mg -da_refine 2
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -pc_type mg -da_refine 3
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -pc_type mg -da_refine 4
```

```
mpiexec -n 12 ./ex19 -ksp_type bcgs -grashof 1e2 -pc_type mg -da_refine 4
lid velocity = 100., prandtl # = 1., grashof # = 100.
 0 SNES Function norm 1.545962539057e+03
  Linear solve converged due to CONVERGED_RTOL iterations 6
 1 SNES Function norm 9.778196290981e+02
  Linear solve converged due to CONVERGED_RTOL iterations 6
 2 SNES Function norm 6.609659458090e+02
  Linear solve converged due to CONVERGED_RTOL iterations 7
 3 SNES Function norm 2.791922927549e+00
  Linear solve converged due to CONVERGED_RTOL iterations 6
 4 SNES Function norm 4.973591997243e-03
  Linear solve converged due to CONVERGED_RTOL iterations 6
 5 SNES Function norm 3.241555827567e-05
  Linear solve converged due to CONVERGED_RTOL iterations 9
 6 SNES Function norm 9.883136583477e-10
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 6
```

Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4
```

Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4
```

```
./ex19 -da_refine 2 -grashof 1.3e4  
lid velocity = 100., prandtl # = 1., grashof # = 13000.  
0 SNES Function norm 7.971152173639e+02  
Linear solve did not converge due to DIVERGED_ITS iterations 10000  
Nonlinear solve did not converge due to DIVERGED_LINEAR_SOLVE iterations 0
```

Oops! Failure in the linear solver? What if we use a stronger preconditioner?

Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4  
./ex19 -da_refine 2 -grashof 1.3e4 -pc_type mg
```

Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4  
./ex19 -da_refine 2 -grashof 1.3e4 -pc_type mg
```

```
./ex19 -da_refine 2 -grashof 1.3e4 -pc_type mg  
lid velocity = 100., prandtl # = 1., grashof # = 13000.  
...  
4 SNES Function norm 3.209967262833e+02  
Linear solve converged due to CONVERGED_RTOL iterations 9  
5 SNES Function norm 2.121900163587e+02  
Linear solve converged due to CONVERGED_RTOL iterations 9  
6 SNES Function norm 1.139162432910e+01  
Linear solve converged due to CONVERGED_RTOL iterations 8  
7 SNES Function norm 4.048269317796e-01  
Linear solve converged due to CONVERGED_RTOL iterations 8  
8 SNES Function norm 3.264993685206e-04  
Linear solve converged due to CONVERGED_RTOL iterations 8  
9 SNES Function norm 1.154893029612e-08  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 9
```

Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4  
./ex19 -da_refine 2 -grashof 1.3e4 -pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -pc_type mg
```

Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4  
./ex19 -da_refine 2 -grashof 1.3e4 -pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -pc_type mg
```

```
lid velocity = 100., prandtl # = 1., grashof # = 13373.  
...  
48 SNES Function norm 3.124919801005e+02  
   Linear solve converged due to CONVERGED_RTOL iterations 17  
49 SNES Function norm 3.124919800338e+02  
   Linear solve converged due to CONVERGED_RTOL iterations 17  
50 SNES Function norm 3.124919799645e+02  
Nonlinear solve did not converge due to DIVERGED_MAX_IT iterations 50
```


Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4  
./ex19 -da_refine 2 -grashof 1.3e4 -pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -pc_type lu
```

Hands-on: Increasing strength of the nonlinearity

What happens as we increase the nonlinearity by raising the Grashof number?

```
./ex19 -da_refine 2 -grashof 1e2  
./ex19 -da_refine 2 -grashof 1e3  
./ex19 -da_refine 2 -grashof 1e4  
./ex19 -da_refine 2 -grashof 1.3e4  
./ex19 -da_refine 2 -grashof 1.3e4 -pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -pc_type lu
```

```
./ex19 -da_refine 2 -grashof 1.3373e4 -pc_type lu  
...  
48 SNES Function norm 3.193724239842e+02  
   Linear solve converged due to CONVERGED_RTOL iterations 1  
49 SNES Function norm 3.193724232621e+02  
   Linear solve converged due to CONVERGED_RTOL iterations 1  
50 SNES Function norm 3.193724181714e+02  
Nonlinear solve did not converge due to DIVERGED_MAX_IT iterations 50
```

A strong linear solver can't help us here. What now?

Let's try combining Newton's method with one of the other nonlinear solvers we mentioned in the introduction, using PETSc's support for nonlinear composition and preconditioning.

To discuss nonlinear composition and preconditioning, we introduce some definitions and notation.

Our prototypical nonlinear equation is of the form

$$F(x) = b \tag{5}$$

and we define the residual as

$$r(x) = F(x) - b \tag{6}$$

We use the notation $x_{k+1} = \mathcal{M}(F, x_k, b)$ for the action of a nonlinear solver.

Nonlinear Composition: Additive

Nonlinear composition consists of a sequence of two (or more) methods \mathcal{M} and \mathcal{N} , which both provide an approximation solution to $F(x) = b$.

In the linear case, application of a stationary solver by defect correct can be written as

$$x_{k+1} = x_k - P^{-1}(Ax_k - b) \quad (7)$$

where P^{-1} is a linear preconditioner. (Richardson iteration applied to a preconditioned system.)

An additive composition of preconditioners P^{-1} and Q^{-1} with weights α and β may be written as

$$x_{k+1} = x_k - (\alpha P^{-1} + \beta Q^{-1})(Ax_k - b) \quad (8)$$

Analogously, for the nonlinear case, additive composition is

$$x_{k+1} = x_k + \alpha \cdot (\mathcal{M}(F, x_k, b) - x_k) + \beta \cdot (\mathcal{N}(F, x_k, b) - x_k) \quad (9)$$

A multiplicative combination of linear preconditioners may be written as

$$\begin{aligned}x_{k+1/2} &= x_k - P^{-1}(Ax_k - b), \\x_{k+1} &= x_{k+1/2} - Q^{-1}(Ax_{k+1/2} - b),\end{aligned}\tag{10}$$

Analogously, for the nonlinear case

$$x_{k+1} = \mathcal{M}(F, \mathcal{N}(F, x_k, b), b)\tag{11}$$

which simply indicates to update the solution using the current solution and residual with the first solver and then update the solution again using the resulting new solution and new residual with the second solver.

Nonlinear Left Preconditioning

Recall that the stationary iteration for our left-preconditioned linear system is

$$x_{k+1} = x_k - P^{-1}(Ax_k - b) \quad (12)$$

And since $Ax_k - b = r$, for the linear case we can write the action of our solver \mathcal{N} as

$$\mathcal{N}(F, x, b) = x_k - P^{-1}r \quad (13)$$

With slight rearranging, we can express the left-preconditioned residual

$$P^{-1}r = x_k - \mathcal{N}(F, x, b) \quad (14)$$

And generalizing to the nonlinear case, the left preconditioning operation provides a modified residual

$$r_L = x_k - \mathcal{N}(F, x, b) \quad (15)$$

Nonlinear Right Preconditioning

For a right preconditioned linear system $AP^{-1}Px = b$, we solve the systems

$$\begin{aligned} AP^{-1}y &= b \\ x &= P^{-1}y \end{aligned} \tag{16}$$

Analogously, we define the right preconditioning operation in the nonlinear case as

$$\begin{aligned} y &= \mathcal{M}(F(\mathcal{N}(F, \cdot, b)), x_k, b) \\ x &= \mathcal{N}(F, y, b) \end{aligned} \tag{17}$$

(Note: In the linear case the above actually reduces to $A(I - P^{-1}A)y = (I - AP^{-1})b$, but the inner solver is applied before the function evaluation (matrix-vector product in the linear case), so we retain the “right preconditioning” name.)

Nonlinear Composition and Preconditioning

Type	Sym	Statement	Abbreviation
Additive	+	$x + \alpha(\mathcal{M}(\mathcal{F}, x, b) - x) + \beta(\mathcal{N}(\mathcal{F}, x, b) - x)$	$\mathcal{M} + \mathcal{N}$
Multiplicative	*	$\mathcal{M}(\mathcal{F}, \mathcal{N}(\mathcal{F}, x, b), b)$	$\mathcal{M} * \mathcal{N}$
Left Prec.	$-_L$	$\mathcal{M}(x - \mathcal{N}(\mathcal{F}, x, b), x, b)$	$\mathcal{M} -_L \mathcal{N}$
Right Prec.	$-_R$	$\mathcal{M}(\mathcal{F}(\mathcal{N}(\mathcal{F}, x, b)), x, b)$	$\mathcal{M} -_R \mathcal{N}$
Inner Lin. Inv.	\backslash	$y = J(x)^{-1}r(x) = K(J(x), y_0, b)$	$\mathcal{N} \backslash K$

Composing Scalable Nonlinear Algebraic Solvers,
Brune, Knepley, Smith, and Tu, SIAM Review, 2015.

For details on using nonlinear composition and preconditioning, see manual pages for SNESCOMPOSITE and SNESGetNPC().

Hands-on: Nonlinear Richardson Preconditioned with Newton

```
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
newtonls -npc_snes_max_it 4 -npc_pc_type mg
```

Hands-on: Nonlinear Richardson Preconditioned with Newton

```
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
newtonls -npc_snes_max_it 4 -npc_pc_type mg
```

```
lid velocity = 100., prandtl # = 1., grashof # = 13373.  
Nonlinear solve did not converge due to DIVERGED_INNER iterations 0
```

Hands-on: Nonlinear Richardson Preconditioned with Newton

```
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
    newtonls -npc_snes_max_it 4 -npc_pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
    newtonls -npc_snes_max_it 4 -npc_pc_type lu
```

Hands-on: Nonlinear Richardson Preconditioned with Newton

```
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
newtonls -npc_snes_max_it 4 -npc_pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
newtonls -npc_snes_max_it 4 -npc_pc_type lu
```

```
lid velocity = 100., prandtl # = 1., grashof # = 13373.  
0 SNES Function norm 7.987708558131e+02  
1 SNES Function norm 8.467169687854e+02  
2 SNES Function norm 7.300096001529e+02  
3 SNES Function norm 5.587232361127e+02  
4 SNES Function norm 3.071143076019e+03  
5 SNES Function norm 3.347748537471e+02  
6 SNES Function norm 1.383297972324e+01  
7 SNES Function norm 1.209841384629e-02  
8 SNES Function norm 8.660606193428e-09  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 8
```

So nonlinear Richardson preconditioned with Newton has let us go further than Newton alone.

Hands-on: Nonlinear Richardson Preconditioned with Newton

```
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
    newtonls -npc_snes_max_it 4 -npc_pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
    newtonls -npc_snes_max_it 4 -npc_pc_type lu  
./ex19 -da_refine 2 -grashof 1.4e4 -snes_type nrichardson -npc_snes_type newtonls  
    -npc_snes_max_it 4 -npc_pc_type lu
```

Hands-on: Nonlinear Richardson Preconditioned with Newton

```
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
    newtonls -npc_snes_max_it 4 -npc_pc_type mg  
./ex19 -da_refine 2 -grashof 1.3373e4 -snes_type nrichardson -npc_snes_type  
    newtonls -npc_snes_max_it 4 -npc_pc_type lu  
./ex19 -da_refine 2 -grashof 1.4e4 -snes_type nrichardson -npc_snes_type newtonls  
    -npc_snes_max_it 4 -npc_pc_type lu
```

```
lid velocity = 100., prandtl # = 1., grashof # = 14000.  
...  
37 SNES Function norm 5.992348444448e+02  
38 SNES Function norm 5.992348444290e+02  
Nonlinear solve did not converge due to DIVERGED_INNER iterations 38
```

We've hit another barrier. What about switching things up?
Let's try preconditioning Newton with nonlinear Richardson.

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
-snes_max_it 1000
```

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
-snes_max_it 1000
```

```
...  
352 SNES Function norm 2.145588832260e-02  
    Linear solve converged due to CONVERGED_RTOL iterations 7  
353 SNES Function norm 1.288292314235e-05  
    Linear solve converged due to CONVERGED_RTOL iterations 8  
354 SNES Function norm 3.219155715396e-10  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 354
```


Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
      -snes_max_it 1000  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 3
```

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
      -snes_max_it 1000  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 3
```

```
...  
23 SNES Function norm 4.796734188970e+00  
   Linear solve converged due to CONVERGED_RTOL iterations 7  
24 SNES Function norm 2.083806106198e-01  
   Linear solve converged due to CONVERGED_RTOL iterations 8  
25 SNES Function norm 1.368771861149e-04  
   Linear solve converged due to CONVERGED_RTOL iterations 8  
26 SNES Function norm 1.065794992653e-08  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 26
```

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
-snes_max_it 1000  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 3  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 4  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 5  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 6  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 7
```

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
      -snes_max_it 1000  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 3  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 4  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 5  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 6  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 7
```

```
lid velocity = 100., prandtl # = 1., grashof # = 14000.  
0 SNES Function norm 8.016512665033e+02  
  Linear solve converged due to CONVERGED_RTOL iterations 11  
1 SNES Function norm 7.961475922316e+03  
  Linear solve converged due to CONVERGED_RTOL iterations 10  
2 SNES Function norm 3.238304139699e+03  
  Linear solve converged due to CONVERGED_RTOL iterations 10  
3 SNES Function norm 4.425107973263e+02  
  Linear solve converged due to CONVERGED_RTOL iterations 9  
4 SNES Function norm 2.010474128858e+02  
  Linear solve converged due to CONVERGED_RTOL iterations 8  
5 SNES Function norm 2.936958163548e+01  
  Linear solve converged due to CONVERGED_RTOL iterations 8  
6 SNES Function norm 1.183847022611e+00  
  Linear solve converged due to CONVERGED_RTOL iterations 8  
7 SNES Function norm 6.662829301594e-03  
  Linear solve converged due to CONVERGED_RTOL iterations 7  
8 SNES Function norm 6.170083332176e-07  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 8
```

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
      -snes_max_it 1000  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 3  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 4  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 5  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 6  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 7
```

Newton preconditioned with nonlinear Richardson can be pushed quite far! Try

```
./ex19 -da_refine 2 -grashof 1e6 -pc_type lu -npc_snes_type nrichardson -npc_snes_max_it 7  
      -snes_max_it 1000
```

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
      -snes_max_it 1000  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 3  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 4  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 5  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 6  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 7
```

Newton preconditioned with nonlinear Richardson can be pushed quite far! Try

```
./ex19 -da_refine 2 -grashof 1e6 -pc_type lu -npc_snes_type nrichardson -npc_snes_max_it 7  
      -snes_max_it 1000
```

```
lid velocity = 100., prandtl # = 1., grashof # = 1e+06  
...  
69 SNES Function norm 4.241700887134e+00  
   Linear solve converged due to CONVERGED_RTOL iterations 1  
70 SNES Function norm 3.238739735055e+00  
   Linear solve converged due to CONVERGED_RTOL iterations 1  
71 SNES Function norm 1.781881532852e+00  
   Linear solve converged due to CONVERGED_RTOL iterations 1  
72 SNES Function norm 1.677710773493e-05  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 72
```

Hands-on: Newton Preconditioned with Nonlinear Richardson

```
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 1  
-snes_max_it 1000  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 3  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 4  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 5  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 6  
./ex19 -da_refine 2 -grashof 1.4e4 -pc_type mg -npc_snes_type nrichardson -npc_snes_max_it 7
```

Newton preconditioned with nonlinear Richardson can be pushed quite far! Try

```
./ex19 -da_refine 2 -grashof 1e6 -pc_type lu -npc_snes_type nrichardson -npc_snes_max_it 7  
-snes_max_it 1000
```

```
lid velocity = 100., prandtl # = 1., grashof # = 1e+06  
...  
69 SNES Function norm 4.241700887134e+00  
Linear solve converged due to CONVERGED_RTOL iterations 1  
70 SNES Function norm 3.238739735055e+00  
Linear solve converged due to CONVERGED_RTOL iterations 1  
71 SNES Function norm 1.781881532852e+00  
Linear solve converged due to CONVERGED_RTOL iterations 1  
72 SNES Function norm 1.677710773493e-05  
Nonlinear solve converged due to CONVERGED_FNORM_RELATIVE iterations 72
```

PETSc's support for extensive runtime experimentation and composition enables discovery of effective approaches when the best solver cannot be determined *a priori*.

Takeaways

PETSc provides a wide assortment of nonlinear solvers through the (SNES) component

Users can build sophisticated solvers from composable algorithmic components:

- Inner, linear solves can employ full range of solvers and preconditioners provided by PETSc KSP and PC
 - Multigrid solvers particularly important for mesh size-independent convergence
- Composite nonlinear solvers can be built analogously, using building blocks from PETSc SNES

Newton-Krylov dominates, but large design space of “composed” nonlinear solvers to explore:

- Not well-explored theoretically or experimentally (interesting research opportunities!)
- Composed nonlinear solvers can be very powerful, though frustratingly fragile
 - Nonlinear Richardson, Gauss-Seidel, or NGMRES with Newton often improves robustness

Further items to explore include

- Nonlinear domain decomposition (SNESASPIN/SNESNASM) and nonlinear multigrid (or Full Approximation Scheme, SNESFAS) methods
- PETSc TS timesteppers use SNES to solve nonlinear problems at each time step
 - Pseudo-transient continuation (TSPSEUDO) can solve highly nonlinear steady-state problems

Our discussion of nonlinear solver algorithms and how to use them via PETSc is mostly orthogonal to the topic of how to run PETSc solvers on GPUs.

Since computing on GPUs has become so important, however, in the slides that follow, we take a brief look at GPU support in PETSc and how its nonlinear solvers can be executed on GPUs.

How PETSc Uses GPUs: Back-End

- Provides several new implementations of PETSc's Vec (distributed vector) and Mat (distributed matrix) classes which allow data storage and manipulation in device (GPU) memory
- Embeds all Vec (and Mat) objects with the ability to track the state of a second “offloaded” copy of the data, and synchronize these two copies of the data (only) when required (“lazy-mirror” model).
- Because higher-level PETSc objects rely on Vec and Mat operations, execution occurs on GPU when appropriate delegated types for Vec and Mat are chosen.

Host and Device Data

```
struct _p_Vec {  
    ...  
    void          *data;           // host buffer  
    void          *spptr;          // device buffer  
    PetscOffloadMask offloadmask;  // which copies are valid  
};
```

Possible Flag States

```
typedef enum {PETSC_OFFLOAD_UNALLOCATED,  
              PETSC_OFFLOAD_GPU,  
              PETSC_OFFLOAD_CPU,  
              PETSC_OFFLOAD_BOTH} PetscOffloadMask;
```

Using GPU Back-Ends in PETSc

Transparently use GPUs for common matrix and vector operations, via runtime options. Currently CUDA/cuSPARSE, HIP/hipSPARSE, Kokkos, and ViennaCL are supported.

CUDA/cuSPARSE usage:

- CUDA matrix and vector types:
`-mat_type aijcusparses -vec_type cuda`
- GPU-enabled preconditioners:
 - GPU-based ILU: `-pc_type ilu -pc_factor_mat_solver_type cusparses`
 - Jacobi: `-pc_type jacobi`

Because PETSc separates high-level control logic from optimized computational kernels, even very complicated hierarchical/multi-level/domain-decomposed/physics-based solvers can run on different architectures by simply choosing the appropriate back-end at runtime; **re-coding is not needed**.

Hands-on: CPU and GPU, detailed performance logging

Run a large version of the driven cavity problem, using multigrid with GPU and SIMD-friendly Chebyshev-Jacobi smoothing (`-mg_levels_pc_type jacobi`), and collect a breakdown by multigrid level (`-pc_mg_log`), both in plain text and flamegraph stack formats.

We get the best CPU-only performance on polaris using 8 MPI ranks:

```
mpiexec -n 8 ./ex19 -da_refine 9 -pc_type mg -mg_levels_pc_type jacobi -pc_mg_log  
-log_view :log_mg_cpu_n8.txt  
  
mpiexec -n 8 ./ex19 -da_refine 9 -pc_type mg -mg_levels_pc_type jacobi -pc_mg_log  
-log_view :log_mg_cpu_n8.stack:ascii_flamegraph
```

Running on GPUs, we get best performance using only one rank per GPU:

```
mpiexec -n 4 ./ex19 -da_refine 9 -pc_type mg -mg_levels_pc_type jacobi -pc_mg_log  
-dm_vec_type cuda -dm_mat_type aijcusparse -log_view_gpu_time -log_view  
:log_mg_gpu_n4.txt  
  
mpiexec -n 4 ./ex19 -da_refine 9 -pc_type mg -mg_levels_pc_type jacobi -pc_mg_log  
-dm_vec_type cuda -dm_mat_type aijcusparse -log_view_gpu_time -log_view  
:log_mg_gpu_n4.stack:ascii_flamegraph
```

Reading -log_view

- Overall summary:

	Max	Max/Min	Avg	Total
Time (sec):	4.106e+01	1.000	4.106e+01	
Objects:	1.201e+03	1.000	1.201e+03	
Flops:	2.577e+10	1.005	2.569e+10	2.055e+11
Flops/sec:	6.276e+08	1.005	6.255e+08	5.004e+09
MPI Msg Count:	1.102e+04	1.543	9.172e+03	7.338e+04
MPI Msg Len (bytes):	4.471e+07	1.655	3.910e+03	2.869e+08
MPI Reductions:	2.806e+03	1.000		

- Also a summary per stage
- Memory usage per stage (based on when it was allocated)
- Time, messages, reductions, balance, flops per event per stage
- Always send `-log_view` when asking performance questions on mailing list!

PETSc Profiling

Event	Count		Time (sec)		Flop		Mess	AvgLen	Reduct	--- Global ---					--- Stage ----					Total Mflop/
	Max	Ratio	Max	Ratio	Max	Ratio				%T	%F	%M	%L	%R	%T	%F	%M	%L	%R	
--- Event Stage 0: Main Stage																				
BuildTwoSided	86	1.0	2.4533e-02	2.4	0.00e+00	0.0	5.0e+02	4.0e+00	8.6e+01	0	0	1	0	3	0	0	2	0	3	0
BuildTwoSidedF	49	1.0	2.0815e-02	2.5	0.00e+00	0.0	0.0e+00	0.0e+00	4.9e+01	0	0	0	0	2	0	0	0	0	2	0
MatMult	349	1.0	6.0068e+00	1.0	4.25e+09	1.0	6.8e+03	6.5e+03	1.8e+01	14	16	9	16	1	35	41	33	43	1	5642
MatMultTranspose	9	1.0	1.0729e-02	1.0	7.11e+06	1.0	3.5e+02	6.4e+02	9.0e+00	0	0	0	0	0	0	0	2	0	0	5287
MatLUFactorSym	1	1.0	4.8302e-05	1.2	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0
MatLUFactorNum	3	1.0	7.8808e-05	1.2	2.73e+04	1.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	2769
MatCopy	2	1.0	9.6890e-06	1.3	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0
MatConvert	1	1.0	3.8523e-05	1.4	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0
MatAssemblyBegin	50	1.0	2.1382e-02	2.4	0.00e+00	0.0	0.0e+00	0.0e+00	4.9e+01	0	0	0	0	2	0	0	0	0	2	0
MatAssemblyEnd	50	1.0	1.1048e-01	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	7.6e+01	0	0	0	0	3	1	0	0	0	3	0
MatGetRowIJ	1	1.0	4.6280e-06	1.5	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0
MatCreateSubMats	3	1.0	2.6247e-04	1.1	0.00e+00	0.0	0.0e+00	0.0e+00	6.0e+00	0	0	0	0	0	0	0	0	0	0	0
MatGetOrdering	1	1.0	3.4175e-05	1.2	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0
MatFDCColorCreate	10	1.0	9.8993e-03	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	2.0e+01	0	0	0	0	1	0	0	0	0	1	0
MatFDCColorSetUp	10	1.0	1.5871e+00	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	8.0e+02	4	0	0	0	29	9	0	0	0	29	0
MatFDCColorApply	30	1.0	2.8531e+00	1.0	2.29e+09	1.0	1.3e+04	3.8e+03	8.7e+01	7	9	18	17	3	17	22	62	49	3	6390
MatFDCColorFunc	630	1.0	9.7597e-01	1.1	2.09e+09	1.0	1.3e+04	3.8e+03	9.0e+00	2	8	18	17	0	6	20	62	49	0	17090
MatRedundantMat	3	1.0	4.1426e-04	1.1	0.00e+00	0.0	0.0e+00	0.0e+00	6.0e+00	0	0	0	0	0	0	0	0	0	0	0
MatMPIConcateSeq	3	1.0	1.0587e-04	1.4	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0
KSPSetUp	36	1.0	4.7154e+00	1.0	3.21e+09	1.0	5.8e+03	4.2e+03	1.0e+03	11	12	8	8	37	28	31	28	23	37	5420
KSPSolve	3	1.0	3.0258e+01	1.0	2.00e+10	1.0	5.4e+04	3.8e+03	1.4e+02	74	78	73	71	5	177	192	258	199	5	5269
KSPGMRESOrthog	319	1.0	3.9607e+00	1.1	3.10e+09	1.0	0.0e+00	0.0e+00	3.2e+02	9	12	0	0	11	22	30	0	0	11	6243
SNESolve	1	1.0	4.0081e+01	1.0	2.58e+10	1.0	7.3e+04	3.9e+03	2.8e+03	98	100	100	100	98	235	248	354	280	99	5127
SNESetUp	1	1.0	7.4249e-01	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	9.0e+00	2	0	0	0	0	4	0	0	0	0	0
SNESFunctionEval	634	1.0	5.7119e-01	1.1	2.19e+09	1.0	1.2e+02	1.4e+04	1.0e+00	1	9	0	1	0	3	21	1	2	0	30591
SNESJacobianEval	30	1.0	4.4082e+00	1.0	2.29e+09	1.0	1.3e+04	3.8e+03	9.2e+02	11	9	18	17	33	26	22	63	49	33	4136
SNESLineSearch	3	1.0	3.1303e-01	1.0	2.52e+08	1.0	1.2e+02	2.0e+04	1.2e+01	1	1	0	1	0	2	2	1	2	0	6428
DMCoarsen	9	1.0	4.5833e-03	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	1.1e+02	0	0	0	0	4	0	0	0	0	4	0
DMCreateInterp	9	1.0	7.5446e-02	1.0	7.11e+06	1.0	3.5e+02	6.4e+02	2.0e+02	0	0	0	0	7	0	0	2	0	7	752
DMCreateInject	9	1.0	1.0279e-02	1.0	0.00e+00	0.0	2.2e+02	2.9e+04	1.4e+02	0	0	0	2	5	0	0	1	6	5	0
DMCreateMat	10	1.0	9.7872e-01	1.0	0.00e+00	0.0	0.0e+00	0.0e+00	7.0e+01	2	0	0	0	2	6	0	0	0	3	0
[...]																				

PETSc Profiling: Multigrid timings breakdown

--- Event Stage 1: MG Apply

BuildTwoSided	2	1.0	8.1144e-05	1.7	0.00e+00	0.0	5.6e+01	4.0e+00	2.0e+00	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
MatMult	1872	1.0	1.7854e+01	1.0	1.28e+10	1.0	3.7e+04	4.4e+03	0.0e+00	43	50	51	57	0	74	83	71	89	0	5724														
MatMultAdd	468	1.0	5.8300e-01	1.1	3.69e+08	1.0	6.1e+03	1.7e+03	0.0e+00	1	1	8	4	0	2	2	12	6	0	5057														
MatMultTranspose	468	1.0	7.5539e-01	1.4	3.70e+08	1.0	6.1e+03	1.7e+03	0.0e+00	1	1	8	4	0	3	2	12	6	0	3905														
MatSolve	52	1.0	4.5501e-04	1.2	1.96e+05	1.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	3452														
MatResidual	468	1.0	4.7088e+00	1.1	3.29e+09	1.0	9.4e+03	4.4e+03	0.0e+00	11	13	13	14	0	19	21	18	22	0	5565														
KSPSolve	988	1.0	1.8321e+01	1.0	1.13e+10	1.0	3.1e+04	4.0e+03	2.0e+00	44	44	42	43	0	75	74	59	67	100	4935														
SFSetup	2	1.0	1.0429e-04	1.4	0.00e+00	0.0	1.1e+02	1.8e+01	2.0e+00	0	0	0	0	0	0	0	0	0	0	0														
SFPack	2912	1.0	1.1586e-02	1.4	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0														
SFUnpack	2912	1.0	1.4369e-03	3.1	2.44e+05	292.8	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	0	0	0	0	0	0														
VecCopy	1404	1.0	6.1932e-01	1.2	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	1	0	0	0	0	2	0	0	0	0	0														
VecSet	1560	1.0	2.3179e-01	1.4	0.00e+00	0.0	0.0e+00	0.0e+00	0.0e+00	0	0	0	0	0	1	0	0	0	0	0														
VecAYPX	2808	1.0	1.6977e+00	1.3	6.58e+08	1.0	0.0e+00	0.0e+00	0.0e+00	4	3	0	0	0	7	4	0	0	0	3089														
VecAXPBYCZ	936	1.0	9.1505e-01	1.1	8.23e+08	1.0	0.0e+00	0.0e+00	0.0e+00	2	3	0	0	0	4	5	0	0	0	7164														
VecPointwiseMult	1872	1.0	1.8484e+00	1.1	3.29e+08	1.0	0.0e+00	0.0e+00	0.0e+00	4	1	0	0	0	7	2	0	0	0	1419														
VecScatterBegin	2912	1.0	4.0778e-02	1.4	0.00e+00	0.0	5.3e+04	3.5e+03	2.0e+00	0	0	72	64	0	0	0	100	100	100	0														
VecScatterEnd	2912	1.0	1.4528e+00	1.3	2.44e+05	292.8	0.0e+00	0.0e+00	0.0e+00	3	0	0	0	0	5	0	0	0	0	1														
PCApply	1924	1.0	1.8517e+00	1.1	3.29e+08	1.0	3.0e+03	6.2e+01	2.0e+00	4																								



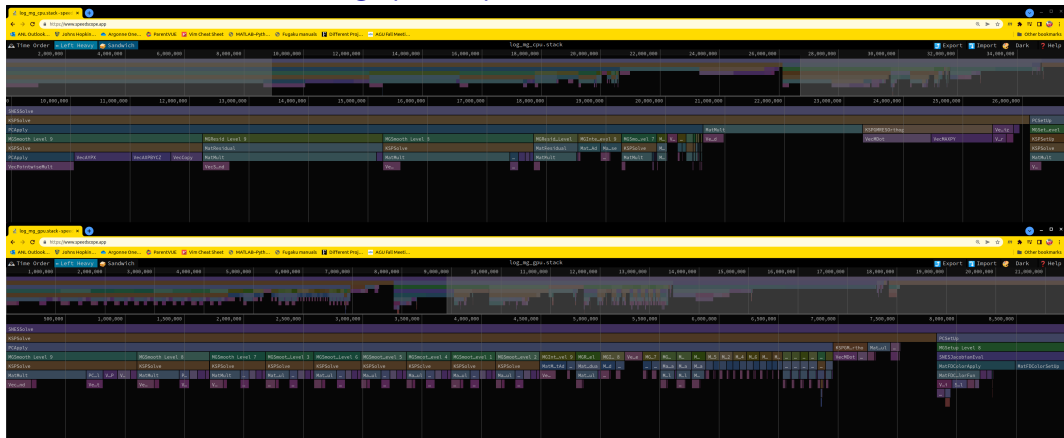
Additional Performance Logging Features: GPU

GPU logging

- When running with GPUs, get additional columns recording GPU flops and data transfers
- Use `-log_view_gpu_time` to get GPU logging for all events (incurs some expense)

```
-----
Event      ... GPU      - CpuToGpu -      - GpuToCpu - GPU
            ... Mflop/s Count   Size   Count   Size  %F
-----
...
KSPSSetUp      ... 82806      112 1.00e+03      35 6.30e+01 100
KSPSolve       ... 119068      65 1.10e-01      65 1.15e-01 100
KSPGMRESOrthog ... 135451      0 0.00e+00      0 0.00e+00 100
SNESSolve      ... 111924     1241 2.75e+03     631 9.30e+02 91
SNESSSetUp     ...      0      0 0.00e+00      0 0.00e+00 0
SNESFunctionEval ...      0      7 3.79e+01      4 3.80e+01 0
SNESJacobianEval ...      0     1022 1.59e+03     527 8.29e+02 0
SNESLineSearch ... 18807      6 5.68e+01      3 2.85e+01 71
...
```


SNES ex19 CPU vs. GPU flame graph comparison



Download the *.stack files to your local machine and then use <https://www.speedscope.app> to generate flame graphs, which will let you examine (interactively) the hierarchy of PETSc events.

One thing to note is the relative distribution of time in MGSmoor steps (part of PCApp1y). See how the steps on the coarse levels all take roughly the same time on the GPU? This points to the high kernel launch latency. What other noteworthy differences can you find?

Further things to try

Speedup (or slowdown!) for GPU vs. CPU

Total time in `SNESolve` tells us the time required to solve our entire problem. Compare these for the CPU and GPU cases to get the overall speedup, but what parts sped up in the GPU case? Which parts actually slowed down? (The slowdowns are mostly due to the fact that the nonlinear function and Jacobian routines in `SNES ex19` do not run on the GPU—verify this by looking at the GPU `%F` column in the text version of the log—and we had to use fewer ranks in this case. See `SNES` tutorial `ex55` in the main development branch of PETSc for an example where these run on the GPU.)

If you'd like to try another GPU-back end, you can try PETSc's Kokkos/Kokkos Kernels one. Run with `-dm_mat_type aijkokkos -dm_vec_type kokkos`.

Experimenting with different multigrid cycle types

Our `SNES ex19` runs defaulted to using multigrid V-cycles. Try running with W-cycles instead by using the option `-pc_mg_cycle_type w`. Unlike V-cycles, W-cycles visit coarse levels many more times than fine ones. What does this do to the time spent in multigrid smoothers for the GPU case vs. the CPU-only one? Should one or the other of these be favored when using GPUs?